

ASSIGNMENT 3

STAT 702 – Data Mining

SUMUKH SAGAR MANJUNATH

1. Based on these data, construct a classification tree for predicting whether an email is “spam” based on other variables. Use 10-fold cross-validation with the 1-SE rule to find the optimal value of the complexity parameter.

```
> library(rpart)
> col_names = read.csv("names.csv",header=F)
> spam = read.csv("spamdata.txt",header=F)
> names(spam) = sapply((1:nrow(col_names)),function(i) toString(col_names[i,1]
]))
> spam$is_spam = factor(spam$is_spam, levels=0:1, labels=c("not_spam", "spam"
))
> is.factor(spam$is_spam)
[1] TRUE
> set.seed(1)
> my.control = rpart.control(xval=10,cp=0)
> cfit = rpart(is_spam ~ ., data=spam, method="class", control = my.control)
> plotcp(cfit)
```

```
> unpruned_spam = printcp(cfit)
```

Classification tree:

```
rpart(formula = is_spam ~ ., data = spam, method = "class", control = my.control)
```

Variables actually used in tree construction:

[1] capital_run_length_average	capital_run_length_longest	capital_run_length_total	char_freq_!
[5] char_freq_\$	char_freq_	char_freq_;	word_freq_1999
[9] word_freq_650	word_freq_address	word_freq_data	word_freq_edu
[13] word_freq_email	word_freq_font	word_freq_free	word_freq_george
[17] word_freq_hp	word_freq_internet	word_freq_money	word_freq_our
[21] word_freq_over	word_freq_re	word_freq_remove	word_freq_technology
[25] word_freq_will	word_freq_you	word_freq_your	

Root node error: 1813/4601 = 0.39404

n= 4601

	CP	nsplit	rel error	xerror	xstd
1	0.47655819	0	1.00000	1.00000	0.018282
2	0.14892443	1	0.52344	0.55378	0.015453
3	0.04302261	2	0.37452	0.45615	0.014366
4	0.03088803	4	0.28847	0.30888	0.012232
5	0.01047987	5	0.25758	0.27910	0.011705
6	0.00827358	6	0.24710	0.26751	0.011489
7	0.00717044	7	0.23883	0.25924	0.011331
8	0.00529509	8	0.23166	0.24986	0.011147
9	0.00441258	14	0.19581	0.23607	0.010867
10	0.00358522	15	0.19140	0.22780	0.010694
11	0.00275786	19	0.17705	0.22339	0.010600
12	0.00257400	22	0.16878	0.21622	0.010445
13	0.00220629	25	0.16106	0.21125	0.010335
14	0.00211436	27	0.15665	0.21125	0.010335
15	0.00165472	33	0.14396	0.21236	0.010360
16	0.00110314	36	0.13900	0.20629	0.010224
17	0.00082736	43	0.13127	0.20243	0.010136
18	0.00055157	47	0.12796	0.20188	0.010124
19	0.00036771	53	0.12466	0.20463	0.010187
20	0.00000000	62	0.12135	0.20574	0.010212

```

> unpruned_spam = as.data.frame(unpruned_spam)
> oneSE_xerr = min(unpruned_spam$xerror) + unpruned_spam$xstd[unpruned_spam$x
error == min(unpruned_spam$xerror)]
> oneSE_xerr
[1] 0.2119991
> optim_cp = max(unpruned_spam$CP[unpruned_spam$xerror < oneSE_xerr])
> optim_cp
[1] 0.002206288

```

What's your estimate of the misclassification rate of the optimal tree? What are the false positive and false negative error rates?

```

> pruned_spam = prune(cfit, cp = optim_cp)
> predicted_spam = predict(pruned_spam, type="vector") - 1
> observed_spam = as.numeric(spam$is_spam) - 1
> missclass = dim(spam[(predicted_spam != observed_spam),,])[1]
> missclass
[1] 292
> total_obs = dim(spam)[1]
> total_obs
[1] 4601
> missclass_rate = (missclass / total_obs)* 100
> missclass_rate
[1] 6.346446
> false_positive = dim(spam[predicted_spam == 1 & observed_spam == 0,])[1]
> false_positive
[1] 105
> false_negative = dim(spam[predicted_spam == 0 & observed_spam == 1,])[1]
> false_negative
[1] 187
> yes_spam = dim(spam[spam$is_spam == "spam",,])[1]
> yes_spam
[1] 1813
> no_spam = dim(spam[spam$is_spam == "not_spam",,])[1]
> no_spam
[1] 2788
> false_positive_rate = false_positive / no_spam
> false_positive_rate
[1] 0.03766141
> false_negative_rate = false_negative / yes_spam
> false_negative_rate
[1] 0.103144

```

How many terminal nodes does your optimal tree have?

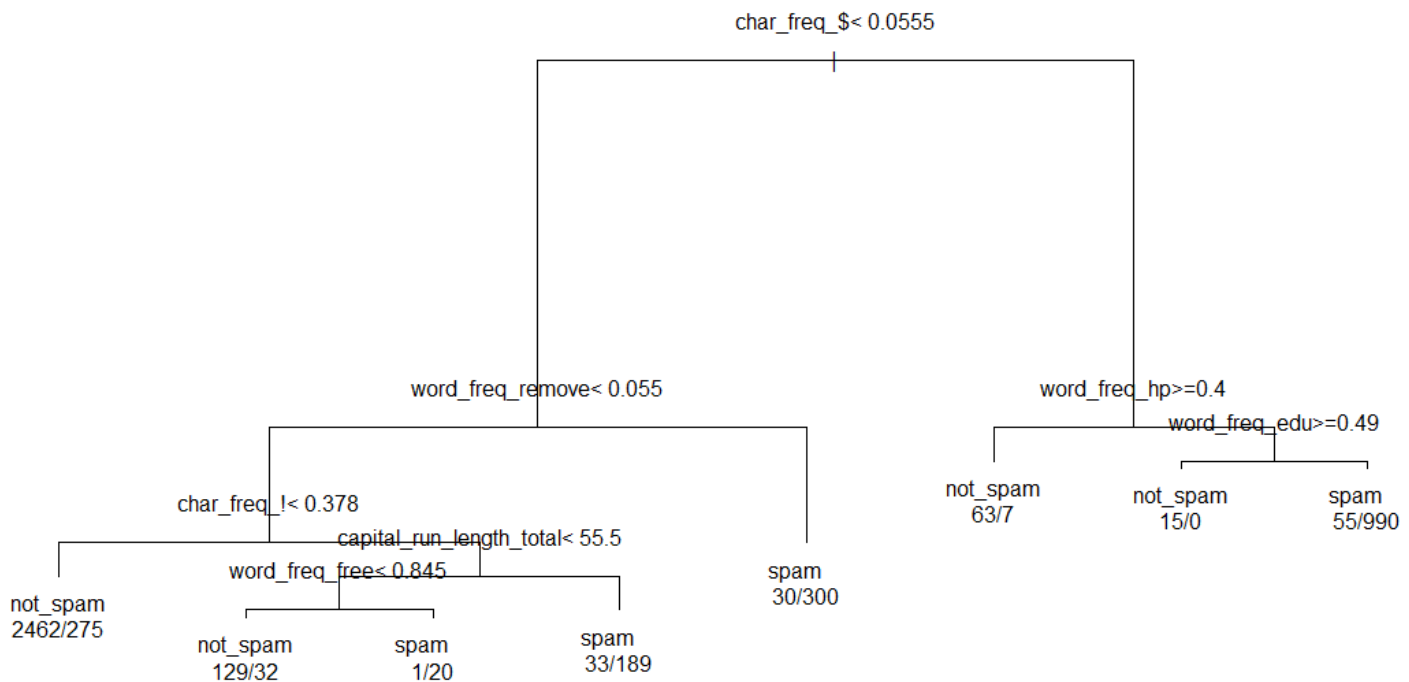
```

> terminal_nodes = unpruned_spam$nsplit[unpruned_spam$CP==optim_cp] + 1
> terminal_nodes
[1] 26

```

Plot the optimal tree. If it is too large, plot a subtree of the optimal tree that has at most 8 terminal nodes.

Pruned subtree with 8 terminal nodes



What are some of the variables that were used in tree construction?

Variables actually used in tree construction:


[1] capital_run_length_average	capital_run_length_longest	capital_run_length_total	char_freq_!
[5] char_freq_\$	char_freq_	char_freq_;	word_freq_1999
[9] word_freq_650	word_freq_address	word_freq_data	word_freq_edu
[13] word_freq_email	word_freq_font	word_freq_free	word_freq_george
[17] word_freq_hp	word_freq_internet	word_freq_money	word_freq_our
[21] word_freq_over	word_freq_re	word_freq_remove	word_freq_technology
[25] word_freq_will	word_freq_you	word_freq_your	

Plot cross validation estimates of errors and training errors of the sequence of pruned trees against the trees' complexity (i.e. number of splits). Compare the two curves (one based on cross validated errors and the other based on training errors.)

```

> splits <- cfit$cptable[, 2]
> terr <- cfit$cptable[, 3]
> xerr <- cfit$cptable[, 4]
> plot(splits, terr, ylim=c(0.1, 1.1), type="l")
> lines(splits, xerr, lty=2)
> title("Cross-validation Error Estimates and Training Error")
  
```

Cross-validation Error Estimates and Training Error



From the plot above, we see that the Cross validation error decreases more drastically than the training error. So, cross validation error is a better choice for pruning.

2. Your classifier in part (1) can be used as a spam filter. One of the possible disadvantages of such a spam filter is that it might filter out too many good (non-spam) emails. Therefore, a better spam filter might be the one which penalizes false positive errors more heavily than false negative errors.

```
> library(rpart)
> col_names = read.csv("names.csv",header=F)
> spam = read.csv("spamdada.txt",header=F)
> names(spam) = sapply((1:nrow(col_names)),function(i) toString(col_names[i,1]
))
> spam$is_spam = factor(spam$is_spam, levels=0:1, labels=c("not_spam", "spam"
))
> is.factor(spam$is_spam)
[1] TRUE
> set.seed(1)
> lmat = matrix(c(0,1,10,0), nrow=2, byrow=F)
> my.control = rpart.control(xval=10,cp=0)
> cfit1 = rpart(is_spam ~ ., data=spam, method="class", control = my.control,
parms=list(loss=lmat))
> unpruned_spam = printcp(cfit1)
> unpruned_spam = as.data.frame(unpruned_spam)
```

```
> unpruned_spam = printcp(cfit1)
```

Classification tree:

```
rpart(formula = is_spam ~ ., data = spam, method = "class", parms = list(loss = lmat),
      control = my.control)
```

Variables actually used in tree construction:

[1]	capital_run_length_average	capital_run_length_longest	capital_run_length_total
[4]	char_freq_!	char_freq_#	char_freq_\$
[7]	char_freq_(char_freq_;	word_freq_000
[10]	word_freq_650	word_freq_address	word_freq_all
[13]	word_freq_edu	word_freq_font	word_freq_free
[16]	word_freq_george	word_freq_hp	word_freq_internet
[19]	word_freq_money	word_freq_order	word_freq_our
[22]	word_freq_people	word_freq_project	word_freq_re
[25]	word_freq_remove	word_freq_technology	word_freq_you
[28]	word_freq_your		

Root node error: 1813/4601 = 0.39404

Root node error: 1813/4601 = 0.39404

n= 4601


	CP	nsplit	rel	error	xerror	xstd
1	0.18422504	0	1.00000	10.0000	0.18282	
2	0.08825152	1	0.81577	5.8097	0.15682	
3	0.05681191	2	0.72752	5.9801	0.15851	
4	0.02647546	4	0.61390	4.6056	0.14382	
5	0.02619967	5	0.58742	3.9691	0.13543	
6	0.02482074	7	0.53502	4.0072	0.13598	
7	0.01765030	9	0.48538	3.4732	0.12809	
8	0.01599559	10	0.46773	3.4997	0.12852	
9	0.01434087	11	0.45174	3.4280	0.12740	
10	0.01185880	12	0.43740	3.3309	0.12581	
11	0.01047987	14	0.41368	3.2537	0.12455	
12	0.00992830	16	0.39272	3.1726	0.12316	
13	0.00827358	17	0.38279	3.2173	0.12390	
14	0.00717044	20	0.35797	3.1197	0.12222	
15	0.00606729	21	0.35080	3.0116	0.12031	
16	0.00551572	22	0.34473	2.9746	0.11963	
17	0.00496415	25	0.32763	2.9371	0.11894	
18	0.00468836	28	0.31274	2.8935	0.11815	
19	0.00441258	32	0.28682	2.9432	0.11904	
20	0.00386100	34	0.27799	2.8742	0.11774	
21	0.00330943	36	0.27027	2.8114	0.11652	
22	0.00220629	42	0.24600	2.7954	0.11622	
23	0.00165472	43	0.24379	2.8229	0.11673	
24	0.00091929	45	0.24049	2.7639	0.11560	
25	0.00000000	52	0.23331	2.7402	0.11507	

```
> oneSE_xerr = min(unpruned_spam$xerror) + unpruned_spam$xstd[unpruned_spam$xerror == min(unpruned_spam$xerror)]
```

```
> optim_cp = max(unpruned_spam$CP[unpruned_spam$xerror < oneSE_xerr])
```

```
> optim_cp
```

```
[1] 0.003309432
```



What's your estimate of the misclassification rate of the optimal tree? What are the false positive and false negative error rates?

```
> pruned_spam = prune(cfit1, cp = optim_cp)
> predicted_spam = predict(pruned_spam, type="vector") - 1
> observed_spam = as.numeric(spam$is_spam) - 1
> missclass = dim(spam[(predicted_spam != observed_spam),,])[1]
> missclass
[1] 445
> total_obs = dim(spam)[1]
> total_obs
[1] 4601
> missclass_rate = (missclass / total_obs)* 100
> missclass_rate
[1] 9.67181
> false_positive = dim(spam[predicted_spam == 1 & observed_spam == 0,])[1]
> false_positive
[1] 5
> false_negative = dim(spam[predicted_spam == 0 & observed_spam == 1,])[1]
> false_negative
[1] 440
> yes_spam = dim(spam[spam$is_spam == "spam",,])[1]
> yes_spam
[1] 1813
> no_spam = dim(spam[spam$is_spam == "not_spam",,])[1]
> no_spam
[1] 2788
> false_positive_rate = false_positive / no_spam
> false_positive_rate
[1] 0.0017934
```

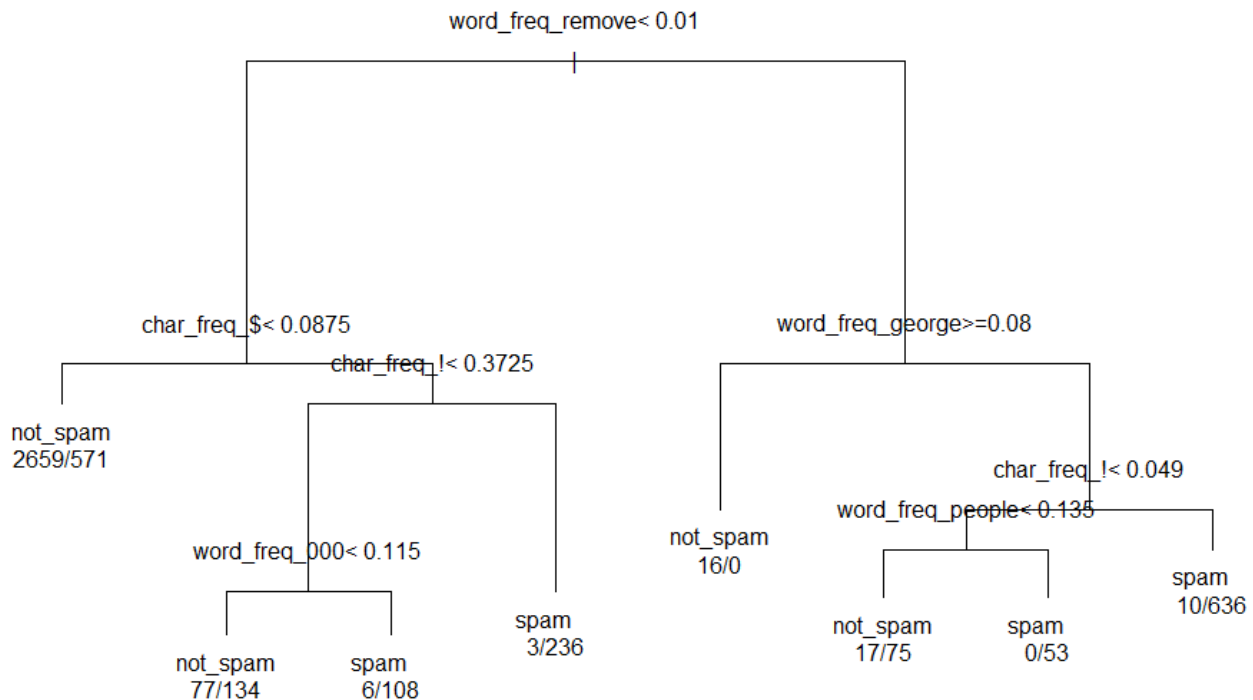
```
> false_negative_rate = false_negative / yes_spam
> false_negative_rate
[1] 0.2426917
```

How many terminal nodes does your optimal tree have?

```
> terminal_nodes = unpruned_spam$nsplit[unpruned_spam$CP==optim_cp] + 1
> terminal_nodes
[1] 37
```

Plot the optimal tree. If it is too large, plot a subtree of the optimal tree that has at most 8 terminal nodes.

Pruned subtree with 8 terminal nodes




```
> pruned_at_8 = prune(cfit1, cp=unpruned_spam$CP[unpruned_spam$nsplit==7])
> plot(pruned_at_8, margin=0.1)
> text(pruned_at_8, use.n=T)
> title("Pruned subtree with 8 terminal nodes")
```

What are some of the variables that were used in tree construction?

variables actually used in tree construction:

[1] capital_run_length_average	capital_run_length_longest	capital_run_length_total	char_freq_!
[5] char_freq_#	char_freq_\$	char_freq_()	char_freq_;
[9] word_freq_000	word_freq_650	word_freq_address	word_freq_all
[13] word_freq_edu	word_freq_font	word_freq_free	word_freq_george
[17] word_freq_hp	word_freq_internet	word_freq_money	word_freq_order
[21] word_freq_our	word_freq_people	word_freq_project	word_freq_re
[25] word_freq_remove	word_freq_technology	word_freq_you	word_freq_your

Plot cross validation estimates of errors and training errors of the sequence of pruned trees against the trees' complexity (i.e. number of splits). Compare the two curves (one based on cross validated errors and the other based on training errors.)



From the plot above, we see that the Cross validation error decreases more drastically than the training error. So, cross validation error is a better choice for pruning.

Which of the two classifiers would you prefer to use as a spam filter and why?

Classifier with the loss matrix is favored because it has lower false positive error meaning chances of falsely classifying a genuine mail as spam is less. This is a desirable expectation from a spam filter since classifying a non-spam mail as spam is more severe than classifying spam mails as non-spam.